The background is a dark navy blue. In the top-left corner, there are several parallel teal lines that form a right-angled corner shape. In the bottom-left corner, there are several parallel teal lines that form a right-angled corner shape. In the bottom-right corner, there are several parallel teal lines that form a right-angled corner shape.

# Machine-Level programming:


## Procedures & Data

# Procedures

- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- Illustration of Recursion

# X86-64 Stack

- ◆ Region of memory managed with stack discipline
- ◆ Grows toward lower addresses
- ◆ Register `%rsp` contains lowest stack address
  - address of “top” element



# X86-64 Stack :Push

- Pushq Src

- Fetch operand at Src
- Decrement %rsp by 8
- Write operand at address given by %rsp

Stack Pointer: %rsp



Stack "Top"

# X86-64 Stack :Pop

- `popq Dest`
  - Read value at address given by `%rsp`
  - Increment `%rsp` by 8
  - Store value at `Dest` (must be register)

Stack Pointer: `%rsp`



Stack "Top"

# Procedures

- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- Illustration of Recursion

# Procedure Control Flow

- Use stack to support procedure call and return
- **Procedure call:** call label
  - Push return address on stack
  - Jump to label
- **Return address:**
  - Address of the next instruction right after call
  - Example from disassembly
- **Procedure Return:** ret
  - Pop address from stack
  - Jump to address

# Control Flow Example #1

## Control Flow Example #1

```
0000000000400540 <multstore>:  
.  
.  
400544: callq 400550 <mult2>  
400549: mov    %rax, (%rbx)  
.  
.
```

```
0000000000400550 <mult2>:  
400550: mov    %rdi,%rax  
.  
.  
400557: retq
```

0x130

0x128

0x120

%rsp

0x120

%rip

0x400544






# Control Flow Example #2

## Control Flow Example #2

```
0000000000400540 <multstore>:  
.  
.  
400544: callq 400550 <mult2>  
400549: mov %rax, (%rbx)  
.  
.
```

```
0000000000400550 <mult2>:  
400550: mov %rdi,%rax  
.  
.  
400557: retq
```




# Control Flow Example #3

## Control Flow Example #3

```
0000000000400540 <multstore>:  
.  
.  
400544: callq 400550 <mult2>  
400549: mov   %rax, (%rbx)  
.  
.
```

```
0000000000400550 <mult2>:  
400550: mov   %rdi,%rax  
.  
.  
400557: retq
```



# Control Flow Example #4

## Control Flow Example #4

```
0000000000400540 <multstore>:  
.  
.  
400544: callq 400550 <mult2>  
400549: mov   %rax, (%rbx)  
.  
.
```

```
0000000000400550 <mult2>:  
400550: mov   %rdi,%rax  
.  
.  
400557: retq
```

0x130

0x128

0x120

%rsp

0x120

%rip

0x400549

.

.

.

# Procedure Data Flow

## Registers

- First 6 arguments

%rdi
%rsi
%rdx
%rcx
%r8
%r9

- Return Value

%rax
------

## Stack

- First 6 arguments

...
Arg <i>n</i>
...
Arg 8
Arg 7

Only allocate stack space  
when needed

# Data Flow Examples

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

```
0000000000400540 <multstore>:
    # x in %rdi, y in %rsi, dest in %rdx
    ...
400541: mov     %rdx,%rbx        # Save dest
400544: callq   400550 <mult2>    # mult2(x,y)
    # t in %rax
400549: mov     %rax,(%rbx)       # Save at dest
    ...
```

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
0000000000400550 <mult2>:
    # a in %rdi, b in %rsi
400550: mov     %rdi,%rax        # a
400553: imul    %rsi,%rax        # a * b
    # s in %rax
400557: retq                               # Return
```

# Managing Local Data

## Stack-Based Languages

### ■ Recursion

- Code must be “Reentrant”
  - Multiple simultaneous instantiations of single procedure
- Need some place to store state of each instantiation
  - Arguments
  - Local variables
  - Return pointer

# Managing Local Data

## Stack-Based Languages

### ■ Stack discipline

- State for given procedure needed for limited time
  - From when called to when return
- Callee returns before caller does

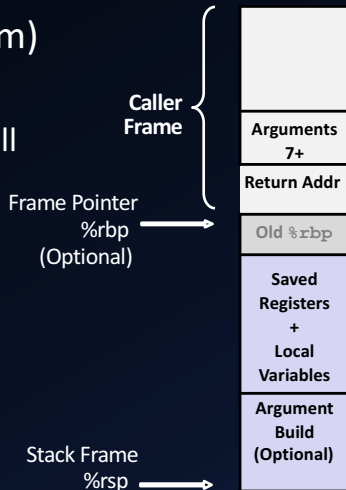
### ■ Stack allocated in Frames

- state for single procedure instantiation

# Stack Frames

## ■ Current Stack Frame(“Top” to Bottom)

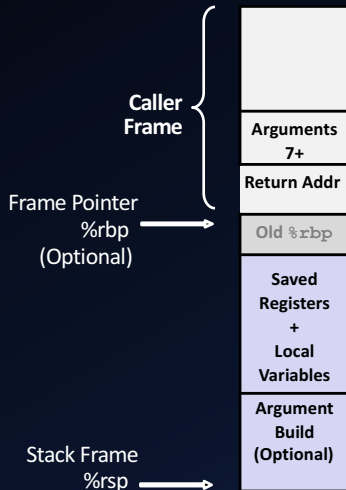
- “Argument build:”  
Parameters for function about to call
- Local variables  
if can't keep in registers
- Saved register context
- Old frame pointer (optional)





# Stack Frames

- Caller Stack Frame
  - Return address
    - Pushed by call instruction
  - Argument for this call



# Register Saving Conventions

- When procedure yoo calls who:
  - yoo is the *caller*
  - who is the *callee*
- Can register be used for temporary storage?

yoo:

```
• • •  
movq $15213, %rdx  
call who  
addq %rdx, %rax  
• • •  
ret
```

who:

```
• • •  
subq $18213, %rdx  
• • •  
ret
```

- Contents of register %rdx overwritten by who
- This could be trouble :something should be done!
  - Need some coordination

# Register Saving Conventions

- When procedure yoo calls who:
  - yoo is the *caller*
  - who is the *callee*
- Can register be used for temporary storage?
- Conventions
  - “*Caller Saved*”
    - Caller saves temporary values in its frame before the call
  - “*callee Saved*”
    - Callee saves temporary values in its frame before using
    - Callee restores them before returning to caller

# X86-64 Linux Register Usage #1

- `%rax`
  - Return value
  - Also caller-saved
  - Can be modified by procedure
- `%rdi,...,%r9`
  - Arguments
  - Also caller-saved
  - can be modified by procedure
- `%r10,%r11`
  - Caller-saved
  - Can be modified by procedure

# X86-64 Linux Register Usage #2

- `%rbx,%r12,%r13,%r14`
  - Callee-saved
  - Callee must save & restore
- `%rdi,...,%r9`
  - Callee-saved
  - Callee must save & restore
  - May be used as frame pointer
  - Can mix & match
- `%rsp`
  - Special form of callee save
  - Restored to original value upon exit from procedure

# Procedures

- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- Illustration of Recursion

# Illustration of Recursion

- Handled Without Special Consideration
  - Stack frames mean that each function call has private storage
    - Saved registers & local variables
    - Saved return pointer
  - Register saving conventions prevent one function call from corrupting another's data
  - Stack discipline follows call / return pattern
    - If P calls Q, then Q returns before P
    - Last-In, First-Out
- Also works for mutual recursion
  - P calls Q; Q calls P

# Data

- Arrays
  - One-dimensional
  - Multi-dimensional(nested)
  - Multi-level
- Structures
  - Allocation
  - Access
  - Alignment
- Floating Point



## ■ Basic Principle

$T$  **A**[ $L$ ];

- Array of data type  $T$  and length  $L$
- Contiguously allocated region of  $L * \text{sizeof}(T)$  bytes in memory
- Identifier **A** can be used as a pointer to array element 0: Type  $T^*$



■ Reference	Type	Value
<code>val[4]</code>	<code>int</code>	3
<code>val</code>	<code>int *</code>	$x$
<code>val+1</code>	<code>int *</code>	$x + 4$
<code>&amp;val[2]</code>	<code>int *</code>	$x + 8$
<code>val[5]</code>	<code>int</code>	??
<code>*(val+1)</code>	<code>int</code>	5 <code>//val[1]</code>
<code>val + i</code>	<code>int *</code>	$x + 4 * i$ <code>//&amp;val[i]</code>

# Multi-dimensional(nested) Arrays

## ■ Declaration

`T A[R][C];`

- 2D array of data type  $T$
- $R$  rows,  $C$  columns
- Type  $T$  element requires  $K$  bytes

## ■ Array Size


- $R * C * K$  bytes

## ■ Arrangement

- Row-Major Ordering

$$\begin{bmatrix} A[0][0] & \cdot & \cdot & \cdot & A[0][C-1] \\ \vdots & & & & \vdots \\ A[R-1][0] & \cdot & \cdot & \cdot & A[R-1][C-1] \end{bmatrix}$$

`int A[R][C];`




# Nested Array Row Access


## ■ Row Vectors

- $A[i]$  is array of  $C$  elements
- Each element of type  $T$  requires  $K$  bytes
- Starting address  $A + i * (C * K)$

```
int A[R][C];
```



# Nested Array Row Access Code



```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
# %rdi = index
leaq (%rdi,%rdi,4),%rax # 5 * index
leaq pgh(,%rax,4),%rax  # pgh + (20 * index)
```

## ■ Row Vector

- `pgh[index]` is array of 5 int's
- Starting address `pgh+20*index`

## ■ Machine Code


- Computes and returns address
- Compute as `pgh + 4*(index+4*index)`

# Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };  
zip_dig mit = { 0, 2, 1, 3, 9 };  
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3  
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
  - 8 bytes
- Each pointer points to array of `int`'s



# Element Access in Multi-Level Array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
ret
```


## ■ Computation

- Element access  $\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{digit}]$
- Must do two memory reads
  - First get pointer to row array
  - Then access element within array

# Array Element Accesses


## Nested array

```
int get_pgh_digit
(size_t index, size_t digit)
{
    return pgh[index][digit];
}
```



## Multi-level array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



Accesses looks similar in C, but address computations very different:

$\text{Mem}[\text{pgh} + 20 * \text{index} + 4 * \text{digit}]$      $\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{digit}]$

# N X N Matrix

## Code

### ■ Fixed dimensions

- Know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele(fix_matrix a,
            size_t i, size_t j)
{
    return a[i][j];
}
```

### ■ Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele(size_t n, int *a,
            size_t i, size_t j)
{
    return a[IDX(n,i,j)];
}
```

### ■ Variable dimensions, implicit indexing

- Now supported by gcc

```
/* Get element a[i][j] */
int var_ele(size_t n, int a[n][n],
            size_t i, size_t j) {
    return a[i][j];
}
```



## n X n Matrix Access

### ■ Array Elements

- Address  $\mathbf{A} + i * (\mathbf{C} * \mathbf{K}) + j * \mathbf{K}$
- $\mathbf{C} = \mathbf{n}, \mathbf{K} = 4$
- Must perform integer multiplication

```
/* Get element a[i][j] */  
int var_ele(size_t n, int a[n][n], size_t i, size_t j)  
{  
    return a[i][j];  
}
```


```
# n in %rdi, a in %rsi, i in %rdx, j in %rcx  
imulq    %rdx, %rdi          # n*i  
leaq     (%rsi,%rdi,4), %rax  # a + 4*n*i  
movl     (%rax,%rcx,4), %eax  # a + 4*n*i + 4*j  
ret
```

# Data

- Arrays
  - One-dimensional
  - Multi-dimensional(nested)
  - Multi-level
- Structures
  - Allocation
  - Access
  - Alignment
- Floating Point

# Structure Representation


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- Structure represented as block of memory
  - Big enough to hold all of the fields
- Fields ordered according to declaration
  - Even if another ordering could yield a more compact representation
- Compiler determines overall size + positions of fields
  - Machine-level program has no understanding of the structures in the source code

## Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



### ■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as  $r + 4 * idx$

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```


```
# r in %rdi, idx in %rsi  
leaq  (%rdi,%rsi,4), %rax  
ret
```

# Following Linked List

## ■ C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```



Element i

Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movslq    16(%rdi), %rax          # i = M[r+16]
    movl      %esi, (%rdi,%rax,4)     # M[r+4*i] = val
    movq      24(%rdi), %rdi         # r = M[r+24]
    testq     %rdi, %rdi              # Test r
    jne       .L11                   # if !=0 goto loop
```

# Alignment Principles

## ■ Aligned Data

- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$
- Required on some machines; advised on x86-64

## ■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
  - Inefficient to load or store datum that spans quad word boundaries
  - Virtual memory trickier when datum spans 2 pages

## ■ Compiler

- Inserts gaps in structure to ensure correct alignment of fields

## Satisfying Alignment with Structures

### ■ Within structure:

- Must satisfy each element's alignment requirement


### ■ Overall structure placement

- Each structure has alignment requirement  $K$ 
  - $K$  = Largest alignment of any element
- Initial address & structure length must be multiples of  $K$

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

### ■ Example:


- $K = 8$ , due to `double` element



## Accessing Array Elements

- Compute array offset  $12 * \text{idx}$ 
  - `sizeof(S3)`, including alignment spacers
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`
  - Resolved during linking

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax # 3*idx  
movzwl a+8(,%rax,4),%eax
```



# Data


- Arrays
  - One-dimensional
  - Multi-dimensional(nested)
  - Multi-level
- Structures
  - Allocation
  - Access
  - Alignment
- Floating Point

# Programming with SSE3

## XMM Registers

■ 16 total, each 16 bytes

■ 16 single-byte integers



■ 8 16-bit integers



■ 4 32-bit integers



■ 4 single-precision floats



■ 2 double-precision floats



■ 1 single-precision float




■ 1 double-precision float



# Scalar & SIMD Operations


## ■ Scalar Operations: Single Precision

`addss %xmm0, %xmm1`




## ■ SIMD Operations: Single Precision

`addps %xmm0, %xmm1`



## ■ Scalar Operations: Double Precision

`addsd %xmm0, %xmm1`



## FP Basics

- Arguments passed in `%xmm0`, `%xmm1`, ...
- Result returned in `%xmm0`
- All XMM registers caller-saved

```
float fadd(float x, float y)
{
    return x + y;
}
```

```
double dadd(double x, double y)
{
    return x + y;
}
```

```
# x in %xmm0, y in %xmm1
addss    %xmm1, %xmm0
ret
```

```
# x in %xmm0, y in %xmm1
addsd    %xmm1, %xmm0
ret
```

## FP Memory Referencing

- Integer (and pointer) arguments passed in regular registers
- FP values passed in XMM registers
- Different mov instructions to move between XMM registers, and between memory and XMM registers

```
double dincr(double *p, double v)
{
    double x = *p;
    *p = x + v;
    return x;
}
```

```
# p in %rdi, v in %xmm0
movapd  %xmm0, %xmm1    # Copy v
movsd   (%rdi), %xmm0    # x = *p
addsd   %xmm0, %xmm1    # t = x + v
movsd   %xmm1, (%rdi)    # *p = t
ret
```

# Other Aspects of FP Code

## ■ Floating-point comparisons

- Instructions `ucomiss` and `ucomisd`
- Set condition codes CF, ZF, and PF

## ■ Using constant values

- Set XMM0 register to 0 with instruction `xorpd %xmm0, %xmm0`
- Others loaded from memory

## 例题3.61

3.61 C编译器为 `var_prod_ele` 产生的代码（图 3-29）不能将它在循环中使用的所有值都放进寄存器中，因此它必须在每次循环时都从存储器中读出  $n$  的值。写出这个函数的 C 代码，使用类似于 GCC 执行的那些优化，但是它的编译代码不会让循环值溢出到存储器中。

回忆一下，处理器只有 6 个寄存器可用来保存临时数据，因为寄存器 `%ebp` 和 `%esp` 不能用于此目的。其中一个寄存器还必须用来保存乘法指令的结果。因此，你必须把循环中的值的数量从 6 个（`result`、`Arow`、`Bcol`、`j`、`n` 和  $4*n$ ）减少到 5 个。

需要找到一个对你那种编译器行之有效的策略。不断尝试各种不同的策略，直到有一种能工作。

3.62 下面的代码计算一个  $n \times n$  矩阵的元素  $i, k$  的乘积。这里  $n$  是一个变量，它的值在编译时是未知的。

```
1  /* Compute i,k of variable matrix product */
2  int var_prod_ele(int n, int A[n][n], int B[n][n], int i, int k) {
3      int j;
4      int result = 0;
5      for (j = 0; j < n; j++)
6          result += A[i][j] * B[j][k];
7
8      return result;
9  }
```

图 3-29 计算变长数组的矩阵乘积的元素  $i, k$ 。编译器执行的优化类似于对定长数组的优化

### 例题3.61

下面是 var\_prod\_ele 循环的汇编代码：

*n stored at %ebp+8*

*Registers: Arow in %esi, Bptr in %ecx, j in %edx,*

*result in %ebx, %edi holds 4\*n*

1	.L30:	loop:
2	movl (%ecx), %eax	Get *Bptr (eax)
3	imull ( <i>Arrow</i> %esi, %edx, 4), %eax	Multiply by Arow[j] <i>Arow</i>
4	addl %eax, %ebx	Add to result <i>result +=</i>
5	addl \$1, %edx	Increment j <i>++j</i>
6	addl %edi, %ecx	Add 4*n to Bptr <i>Bptr</i>
7	cmpl %edx, 8(%ebp)	Compare n:j <i>cmp n:j</i>
8	jg .L30	If >, goto loop



## 例题3.61

我们看到程序既使用了伸缩过的值  $4n$ （寄存器 `%edi`）来增加 `Bptr`，也使用了存储在相对于 `%ebp` 偏移量为 8 处的  $n$  的实际值来检查循环的边界。C 代码中并没有体现出需要这两个值，但是由于指针运算的伸缩，才使用了这两个值。每次循环中，代码从存储器中取出  $n$  的值，检查循环是否终止（第 7 行）。这是一个寄存器溢出（register spilling）的例子：没有足够多的寄存器来保存需要的临时数据，因此编译器必须把一些局部变量放在存储器中。在这个情况下，编译器选择把  $n$  溢出，因为它是一个“只读”的值——在循环中不会改变它的值。因为 IA32 处理器的寄存器数量太少，必须常常将循环值溢出到存储器中。通常，读存储器完成起来比写存储器要容易得多，因此将只读变量溢出是比较合适的。关于如何改进这段代码以避免寄存器溢出，请参见家庭作业 3.61。

3.61

```
int var_prod_ele(int n, int A[n][n], int B[n][n], int i, int k)
{
    int j = n-1;
    int result = 0;
    for(; j!=-1; --j)
        result += A[i][j] * B[j][k];
    return result;
}
```

但是这样得到的结果仍然会使用到存储器。

按下面的代码，循环里面貌似就没有用到存储器。

但是用到了一个常量 4，就是增加 a 的时候，会 add 4。

只需要 result, a, e, b, 4n 这五个变量。

```
int var_prod_ele(int n, int A[n][n], int B[n][n], int i, int k)
{
    int result = 0;
    int *a = &A[i][0];
    int *b = &B[0][k];
    int *e = &A[i][n];
    for(;a!=e;)
    {
        result += *a * *b;
        b+=n;
        a++;
    }
}
```

```
    return result;  
}
```

下面是其汇编代码的循环部分：

edi 是  $4*n$ ，ebx 和 edx 分别是 b 和 a，esi 是 e，eax 是 result。

ecx 是用于存储乘法的寄存器。



L4:

```
movl (%ebx), %ecx  
imull (%edx), %ecx  
addl %ecx, %eax  
addl %edi, %ebx  
addl $4, %edx  
cmpl %edx, %esi  
jneL4
```